

# ECE326

## PROGRAMMING LANGUAGES

### **Lecture 34 : Concurrent Programming**

Kuei (Jack) Sun


ECE

University of Toronto

Fall 2019

# Concurrent Programming

- Multiple tasks execute simultaneously
- Thread
  - Independent sequence of execution
    - Has its own stack, but shares the heap with other threads
- Parallel computing
  - Threads executing at same physical time instant
    - Only possible if each thread runs on its own processor
- Concurrency
  - Threads may *interleave* on the same processor



# Purpose

- Speed up program
  - Usually with more people, a job can get done faster
- Criteria for speed up
  - Threads can work relatively independently
    - Seldom need to wait for other threads
      - E.g. to access shared data
      - E.g. to wait for input produced by another thread
  - Threads are often waiting for IO (e.g. read from disk, network)
    - Only important if threads are sharing a processor
    - While one thread waits, other threads can still do work on processor

# Concurrent Programming

- Most programming languages provide *library support*
  - Creating and managing threads done through function calls

```
use std::thread;
use std::time::Duration;

thread::spawn(|| {
    for i in 1..10 {
        println!("hi number {} from thread!", i);
        thread::sleep(Duration::from_millis(1));
    }
});
```

spawn can take a  
*closure* as argument

- Go has language support for creating threads

```
go f(x, y, z); // starts a new thread (aka goroutine)
```

# Basics

- A program always starts with one thread: main
- main creates new threads, and those can create more
- Creator *should* wait for the threads it created to end

```
fn main() {  
    let handle = thread::spawn(|| {  
        for i in 1..10 {  
            println!("hi number {} from thread!", i);  
        }  
    });  
  
    for i in 1..5 {  
        println!("hi number {} from main!", i);  
    }  
  
    handle.join().unwrap();    // wait for created thread to finish  
}
```

# Ownership

- A thread can potentially live longer than its creator
  - E.g. the creator chooses not to call join before exiting
- Problem arises if closure references outer variable
  - Therefore, all outer variables must be “moved” into closure

```
fn main() {  
    let v = vec![1, 2, 3];  
    let handle = thread::spawn(move || {  
        println!("Here's a vector: {:?}", v);  
    });  
    // you may no longer use 'v' here.  
    // main may exit before thread does!  
}
```

By default, closures  
borrow outer variables

Must specify move to  
make closure move  
outer variables into it.

# Challenge

- Sharing data
  - Ownership
    - Threads need to jointly own an object
  - Updates to same data can result in *race condition*
    - Caused by problematic interleaving of threads
      - Depending on timing of thread execution, which is difficult to control
    - Race condition can lead to unexpected and often incorrect results
- Synchronization
  - Threads may need to communicate with each others
  - One thread may need to wait for another thread to advance

# Reference Counting

- A commonly used technique to share an object
- Analogy
  - First person to walk into living room turns on TV
  - Subsequent people entering can sit down immediately
  - Last person to leave will turn off the TV
- Reference Counting
  - Creator of object sets reference count to 1
    - Others will increment count before use
  - Everyone decrements count after use
    - If count is 0, free the object



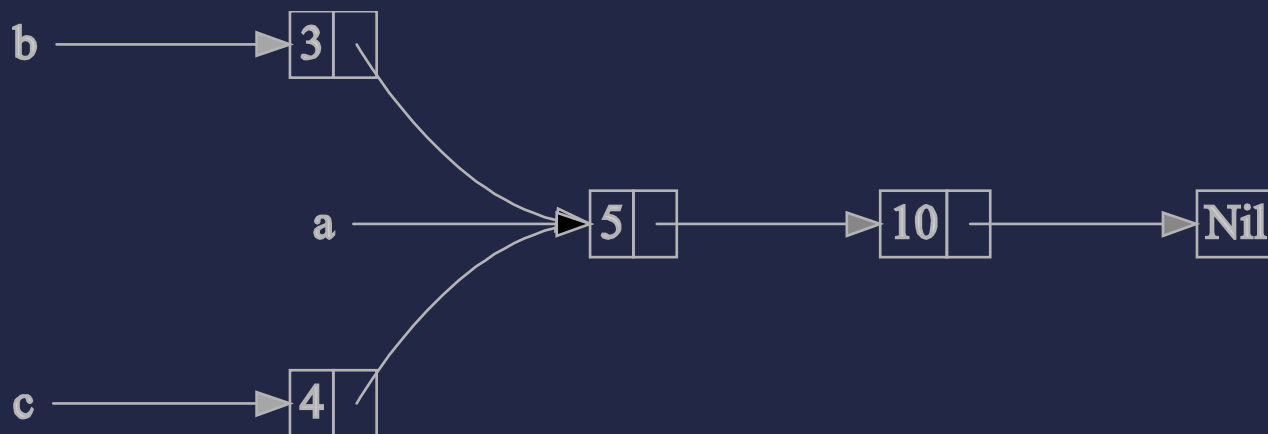
# Smart Pointer

- A wrapper class over a pointer, and acts like a pointer
- C++ Example
  - `unique_ptr`
    - Automatically frees pointed-to object when it goes out of scope
  - `shared_ptr`
    - A reference counting smart pointer
    - Allows multiple threads to share pointed-to object
    - Last reference holder will delete the object
      - May not be the original creator of the object

# Rc<T>

- Allows sharing data in single-threaded setting

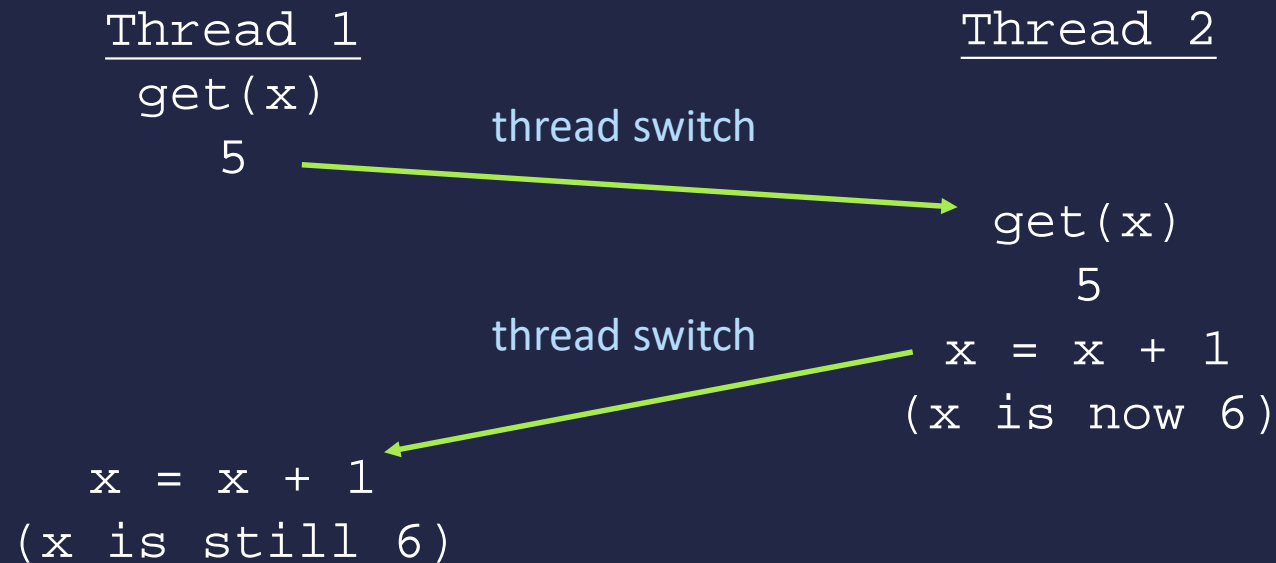
```
enum List { Cons(i32, Rc<List>), Nil, }  
use crate::List::{Cons, Nil};  
use std::rc::Rc;  
  
fn main() {  
    let a = Rc::new(Cons(5, Rc::new(Cons(10, Rc::new(Nil)))));  
    let b = Cons(3, Rc::clone(&a));  
    let c = Cons(4, Rc::clone(&a));  
}
```



Each time Rc::clone is called, reference count of a increases by 1.

# Lost Update

- One potential problem caused by race condition
  - Assume both threads are running on same processor



- Solution: atomic instructions

# Arc<T>

- Allows sharing data across different threads
  - A in Arc stands for *atomic*
- Atomic instruction
  - A single, interruptible instruction on processor
  - Can complete without interference from other threads
  - Generally not used because it takes longer than if split
  - E.g. fetch-and-add

```
function FetchAndAdd(address location, int inc) {  
    int value := *location  
    *location := value + inc  
    return value  
}
```

# Arc<T>

- Arc<T> uses atomic instructions to update counter
  - Unlike regular Rc<T>, which is not *thread safe*
- Thread safety
  - Function that behaves correctly during simultaneous execution by multiple threads
    - E.g. freedom from race condition

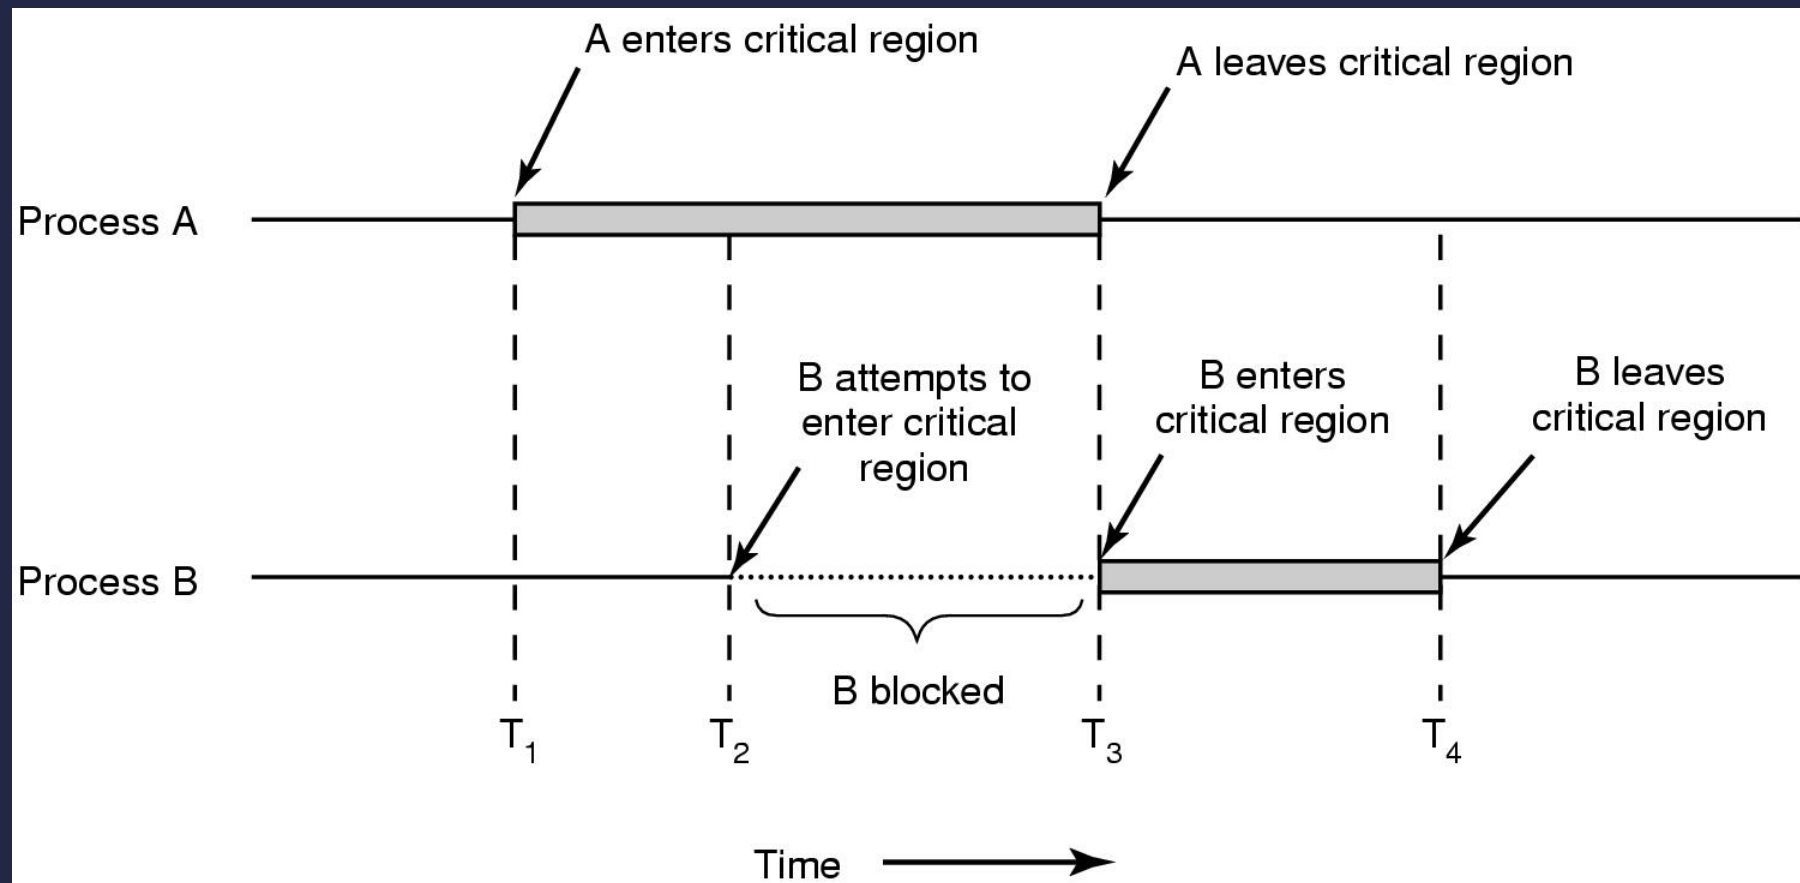
```
use std::sync::Arc;  
let foo = Arc::new(vec![1.0, 2.0, 3.0]);  
// The two syntaxes below are equivalent.  
let a = foo.clone();  
let b = Arc::clone(&foo);  
// a, b, and foo all point to the same shared memory location
```

# Arc<T>

- Limitations
  - Arc<T> makes sharing objects thread safe
  - However, it does *not* make using the objects thread safe
    - E.g. methods of the object may be thread unsafe
  - Object within Arc<T> is *immutable*
- Need
  - A construct that allows shared mutable object
  - A construct that makes using objects thread safe

# Mutual Exclusion

- Ensures only one thread can access shared data at once



# Mutex<T>

- Provides mutual exclusion
- When mutex is locked, no other thread can use object
  - Locking mutex creates a MutexGuard

```
use std::sync::Mutex;
```

```
fn main() {  
    let m = Mutex::new(5);  
    {  
        // num is a MutexGuard around the data  
        let mut num = m.lock().unwrap();  
        *num = 6;  
    } // num goes out of scope and unlocks m  
  
    println!("m = {:?}", m);  
}
```



# Mutex<T>

- Provides *interior mutability*
  - The mutex is immutable, but the data it contains is mutable
- Caveat
  - Mutex is not sharable (one ownership rule)
- Must be combined with Arc<T>

```
let counter = Arc::new(Mutex::new(0));
for _ in 0..10 {
    let counter = Arc::clone(&counter);
    let handle = thread::spawn(move || {
        let mut num = counter.lock().unwrap();
        *num += 1;
    });
}
```

# synchronized

- Language support in Java for mutual exclusion

```
class ThreadedSend extends Thread {  
    private String msg;  
    Sender sender;          // shared among different threads  
  
    ThreadedSend(String m, Sender obj) {  
        msg = m; sender = obj;  
    }  
  
    public void run() {  
        // Only one thread can send message at a time.  
        synchronized(sender) {  
            // synchronizing the send object  
            sender.send(msg);  
        }  
    }  
}
```

# synchronized

- Alternatively, can make an entire method critical region

```
class Sender {  
    // Same effect as previous slide, only one thread can send  
    public synchronized void send(String msg) {  
        System.out.println("Sending\t" + msg );  
        try {  
            Thread.sleep(1000);  
        }  
        catch (Exception e) {  
            System.out.println("Thread interrupted.");  
        }  
        System.out.println("\n" + msg + "Sent");  
    }  
}
```

# Message Passing

- Threads communicate by sending message with data
- Allows threads to *synchronize*
  - i.e. thread waits for condition to satisfy before continuing
- Thread sleeps while waiting for message
- Another thread can wake it up by sending a message
- Bad alternative: Polling (or busy looping)
  - Continuously lock shared data to check on condition in a loop
  - Reduces performance of entire system

# channel<T>

- Creates a sender and a receiver end, thread safe
- Must send/receive same data type

```
use std::thread;
use std::sync::mpsc;

fn main() {
    let (tx, rx) = mpsc::channel();
    thread::spawn(move || {
        let val = String::from("hi");
        tx.send(val).unwrap(); // val moves into send()
    });
    let received = rx.recv().unwrap();
    println!("Got: {}", received);
}
```

# mpsc::channel

- Multiple producer, single consumer
- iteration on rx finishes when channel is closed
  - i.e. when all senders close their end

```
let (tx, rx) = mpsc::channel();  
for i in 1..10 {  
    let tx = mpsc::Sender::clone(&tx);  
    thread::spawn(move || {  
        tx.send(String::from("hello")).unwrap();  
    });  
}  
for received in rx {  
    println!("Got: {}", received);  
}
```

Can clone tx to allow  
for multiple producers.  
rx cannot be cloned!